

DSC550 Week 4

April 7, 2024

```
[30]: import pandas as pd
file_path = '/Users/nickblackford/Desktop/Python/auto-mpg.csv'
df = pd.read_csv(file_path)
df.head()
```

```
[30]:      mpg  cylinders  displacement  horsepower  weight  acceleration  model year  \
0   18.0          8         307.0          130    3504           12.0         70
1   15.0          8         350.0          165    3693           11.5         70
2   18.0          8         318.0          150    3436           11.0         70
3   16.0          8         304.0          150    3433           12.0         70
4   17.0          8         302.0          140    3449           10.5         70

      origin          car name
0         1  chevrolet chevelle malibu
1         1      buick skylark 320
2         1    plymouth satellite
3         1      amc rebel sst
4         1      ford torino
```

```
[31]: # remove car name
df = df.drop(columns=['car name'])
```

```
[32]: # check horsepower type
print(df['horsepower'].dtype)

# horsepower is likely type 'string' because it includes 'NA values'
```

object

```
[33]: # Convert horsepower to numeric
df['horsepower'] = pd.to_numeric(df['horsepower'], errors='coerce')

# Calculate the mean of the horsepower
mean_horsepower = df['horsepower'].mean()

# Replace NA values with mean
df['horsepower'].fillna(mean_horsepower, inplace=True)
```

```
[34]: # Create dummy variables for the 'origin' column
df = pd.get_dummies(df, columns=['origin'], prefix='origin')
```

```
[35]: df.head()
```

```
[35]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	\
0	18.0	8	307.0	130.0	3504	12.0	
1	15.0	8	350.0	165.0	3693	11.5	
2	18.0	8	318.0	150.0	3436	11.0	
3	16.0	8	304.0	150.0	3433	12.0	
4	17.0	8	302.0	140.0	3449	10.5	

	model	year	origin_1	origin_2	origin_3
0		70	True	False	False
1		70	True	False	False
2		70	True	False	False
3		70	True	False	False
4		70	True	False	False


```
[37]: # Calculate the correlation matrix
corr_matrix = df.corr()
corr_matrix
```

```
[37]:
```

	mpg	cylinders	displacement	horsepower	weight	\
mpg	1.000000	-0.775396	-0.804203	-0.771437	-0.831741	
cylinders	-0.775396	1.000000	0.950721	0.838939	0.896017	
displacement	-0.804203	0.950721	1.000000	0.893646	0.932824	
horsepower	-0.771437	0.838939	0.893646	1.000000	0.860574	
weight	-0.831741	0.896017	0.932824	0.860574	1.000000	
acceleration	0.420289	-0.505419	-0.543684	-0.684259	-0.417457	
model year	0.579267	-0.348746	-0.370164	-0.411651	-0.306564	
origin_1	-0.568192	0.604351	0.651407	0.486083	0.598398	
origin_2	0.259022	-0.352861	-0.373886	-0.281258	-0.298843	
origin_3	0.442174	-0.396479	-0.433505	-0.321325	-0.440817	


	acceleration	model year	origin_1	origin_2	origin_3
mpg	0.420289	0.579267	-0.568192	0.259022	0.442174
cylinders	-0.505419	-0.348746	0.604351	-0.352861	-0.396479
displacement	-0.543684	-0.370164	0.651407	-0.373886	-0.433505
horsepower	-0.684259	-0.411651	0.486083	-0.281258	-0.321325
weight	-0.417457	-0.306564	0.598398	-0.298843	-0.440817
acceleration	1.000000	0.288137	-0.250806	0.204473	0.109144
model year	0.288137	1.000000	-0.139883	-0.024489	0.193101
origin_1	-0.250806	-0.139883	1.000000	-0.597198	-0.643317
origin_2	0.204473	-0.024489	-0.597198	1.000000	-0.229895
origin_3	0.109144	0.193101	-0.643317	-0.229895	1.000000

```
[39]: import seaborn as sns
import matplotlib.pyplot as plt
# Plot the heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, annot=True, fmt=".2f", cmap="coolwarm")
plt.title("Correlation Matrix of Auto MPG Dataset")
plt.show()
```



Based on the heatmap, cylinders, displacement, horsepower, and weight are all strongly, negatively correlated with mpg

```
[40]: # Plot mpg versus weight
plt.scatter(df['weight'], df['mpg'])
plt.xlabel('Weight')
plt.ylabel('MPG')
plt.title('MPG vs. Weight')
plt.show()
```



The scatter plot has a clear downward trend, representing the negative correlation between mpg and weight.

```
[44]: from sklearn.model_selection import train_test_split

# Define the features and the target
X = df.drop(columns=['mpg'])
y = df['mpg']

# Split the data into 80% training data and 20% test data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

# Check the size of the training and test sets
print(len(X_train))
print(len(X_test))
```

318
80

```
[45]: from sklearn.linear_model import LinearRegression
```

```
# Initialize the linear regression model
```

```
model = LinearRegression()
```

```
# Train the model on the training data
```

```
model.fit(X_train, y_train)
```

```
[45]: LinearRegression()
```

```
[46]: from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error
import numpy as np
```

```
# Make predictions on the training set
```

```
y_train_pred = model.predict(X_train)
```

```
# Calculate R2, RMSE, and MAE for the training set
```

```
r2_train = r2_score(y_train, y_train_pred)
```

```
rmse_train = np.sqrt(mean_squared_error(y_train, y_train_pred))
```

```
mae_train = mean_absolute_error(y_train, y_train_pred)
```

```
# Make predictions on the test set
```

```
y_test_pred = model.predict(X_test)
```

```
# Calculate R2, RMSE, and MAE for the test set
```

```
r2_test = r2_score(y_test, y_test_pred)
```

```
rmse_test = np.sqrt(mean_squared_error(y_test, y_test_pred))
```

```
mae_test = mean_absolute_error(y_test, y_test_pred)
```

```
# Print the results
```

```
print("Training set metrics:")
```

```
print(f"R2: {r2_train:.2f}")
```

```
print(f"RMSE: {rmse_train:.2f}")
```

```
print(f"MAE: {mae_train:.2f}")
```

```
print("\nTest set metrics:")
```

```
print(f"R2: {r2_test:.2f}")
```

```
print(f"RMSE: {rmse_test:.2f}")
```

```
print(f"MAE: {mae_test:.2f}")
```

Training set metrics:

R2: 0.82

RMSE: 3.37

MAE: 2.61

Test set metrics:

R2: 0.84

RMSE: 2.89

MAE: 2.29

Interpretation:

R2: Our model explains 84% of the variability in mpg.

RMSE: The average error between the actual mpg and our model's predicted mpg is 2.89 mpg

MAE: The average magnitude between our model and actual mpg is 2.29 mpg.

```
[57]: # Try elastic net regression to address multicollinearity concerns
from sklearn.linear_model import ElasticNet

elastic_net_model = ElasticNet(alpha=0.1, l1_ratio=0.5)
elastic_net_model.fit(X_train, y_train)
```

```
[57]: ElasticNet(alpha=0.1)
```

```
[58]: # Make predictions on the training set
y_train_pred_en = elastic_net_model.predict(X_train)

# Calculate R2, RMSE, and MAE for the training set
r2_train_en = r2_score(y_train, y_train_pred_en)
rmse_train_en = np.sqrt(mean_squared_error(y_train, y_train_pred_en))
mae_train_en = mean_absolute_error(y_train, y_train_pred_en)

# Make predictions on the test set
y_test_pred_en = elastic_net_model.predict(X_test)

# Calculate R2, RMSE, and MAE for the test set
r2_test_en = r2_score(y_test, y_test_pred_en)
rmse_test_en = np.sqrt(mean_squared_error(y_test, y_test_pred_en))
mae_test_en = mean_absolute_error(y_test, y_test_pred_en)

# Print the results
print("Elastic Net Regression Metrics:")
print("Training set:")
print(f"R2: {r2_train_en:.2f}")
print(f"RMSE: {rmse_train_en:.2f}")
print(f"MAE: {mae_train_en:.2f}")

print("\nTest set:")
print(f"R2: {r2_test_en:.2f}")
print(f"RMSE: {rmse_test_en:.2f}")
print(f"MAE: {mae_test_en:.2f}")
```

Elastic Net Regression Metrics:

Training set:

R2: 0.82

RMSE: 3.39

MAE: 2.61

Test set:

R2: 0.84

RMSE: 2.92

MAE: 2.32

Interpretation:

R2: Our model explains 84% of the variability in mpg.

RMSE: The average error between the actual mpg and our model's predicted mpg is 2.92 mpg

MAE: The average magnitude between our model and actual mpg is 2.32 mpg.

The elastic net actually did slightly worse than our standard linear regression model.

[]: